



Project 3: CS61CPU

Part A Deadline: Monday, October 25, 11:59:59 PM PT

Part B Deadline: Wednesday, November 10, 11:59:59 PM PT

In this project, you will be building a CPU that runs actual RISC-V instructions.

Setup: Git

1. Visit [Galloc](#). Log in and start the Project 3 assignment. This will create a GitHub repository for your work.
2. Clone the repository on your workspace. We recommend using your local machine. Windows users should clone outside WSL (Git Bash is recommended).

```
$ git clone https://github.com/61c-student/fa21-proj3-USERNAME.git proj3-cs61cpu
```

(replace **USERNAME** with your GitHub username)

3. Navigate to your repository:

```
$ cd proj3-cs61cpu
```

4. Add the starter repository as a remote:

```
$ git remote add starter https://github.com/61c-teach/fa21-proj3-starter.git
```

Setup: Logisim

This project is done in Logisim. To get started, navigate to the **proj3-cs61cpu** directory and run **java -jar tools/logisim-evolution.jar**. This should open the Logisim graphical interface.

☐ Dark Mode



Part A: addi

In this part, you will design a skeleton CPU that can execute the `addi` instruction.

Task 1: Arithmetic Logic Unit (ALU)

Fill in the ALU in `alu.circ` so that it can perform the required arithmetic calculations.

Input Name	Bit Width	Description
A	32	Data to use for Input A in the ALU operation
B	32	Data to use for Input B in the ALU operation
ALUSel	4	Selects which operation the ALU should perform (see the list of operations with corresponding switch values below)

Output Name	Bit Width	Description
Result	32	Result of the ALU operation

Below is the list of ALU operations for you to implement, along with their associated ALUSel values. `add` is already made for you. You are allowed and encouraged to use built-in Logisim components to implement the arithmetic operations.

ALUSel Value	Instruction
0	add: <code>Result = A + B</code>
1	sll: <code>Result = A << B</code>
2	slt: <code>Result = (A < B (signed)) ? 1 : 0</code>
3	Unused
4	xor: <code>Result = A ^ B</code>

ALUSel Value	Instruction
5	srl: $\text{Result} = (\text{unsigned}) A \gg B$
6	or: $\text{Result} = A \mid B$
7	and: $\text{Result} = A \& B$
8	mul: $\text{Result} = (\text{signed}) (A * B)[31:0]$
9	mulh: $\text{Result} = (\text{signed}) (A * B)[63:32]$
10	Unused
11	mulhu: $\text{Result} = (A * B)[63:32]$
12	sub: $\text{Result} = A - B$
13	sra: $\text{Result} = (\text{signed}) A \gg B$
14	Unused
15	bsel: $\text{Result} = B$

Some additional tips:

- You might find bit splitters or extenders useful when implementing shift operations.
- The result of multiplying 2 32-bit numbers can be up to 64 bits of information, but we're limited to 32-bit data lines, so `mulh` and `mulhu` are used to get the upper 32 bits of the product. The `Multiplier` component has a `Carry Out` output, with the description: "the upper bits of the product". This might be particularly useful for certain multiply operations.
- The comparator component might be useful for implementing instructions that involve comparing inputs.
- A multiplexer (MUX) might be useful when deciding between operation outputs. In other words, consider simply processing the input for all operations, and then outputting the one of your choice.
- The ALU tests for Part A only use ALUSel values for defined instructions, so your design doesn't need to worry about the unused values.

Testing

Here's a [companion video](#) that goes with this section.

We've provided some tests for each task, located in subdirectories under `tests/`. For example, the ALU tests are in `tests/part-a/alu/`. When these tests are run, the outputs from your circuits are saved in a `tests/part-a/alu/student-output/` subdirectory.

For example, to run the ALU tests:

```
$ python3 test.py tests/part-a/alu/
```

You can also specify a single test circuit, or a grandparent/great-grandparent directory:

```
$ python3 test.py tests/part-a/alu/alu-add.circ
$ python3 test.py tests/part-a/
$ python3 test.py tests/
```

After the tests finish running, your ALU circuit's outputs will be saved under `tests/part-a/alu/student-output/` with a `-student.out` suffix (e.g. `alu-add-student.out`). The corresponding reference outputs can be found at `tests/part-a/alu/reference-output/` with a `-ref.out` suffix (e.g. `alu-add-ref.out`).

We've also provided `format-output.py`, which accepts a path to an output file and displays the output in a more readable format (left-aligned **hexadecimal** numbers). For example, to get the reference output of the `alu-add` sanity test in readable format, you would do:

```
$ python3 tools/format-output.py tests/part-a/alu/reference-output/alu-add-ref.out
```

If you want to see the difference between your output and the reference solution, you can put the readable outputs into temporary files and `diff` them. For example, for the `alu-add` test, you would do:

```
$ python3 tools/format-output.py tests/part-a/alu/reference-output/alu-add-ref.out > reference.out
$ python3 tools/format-output.py tests/part-a/alu/student-output/alu-add-student.out > student.out
$ diff reference.out student.out
```

Note: If the lines are wrapping, try resizing your terminal window (or try a slightly smaller font). Or see the following note.

Experimental Note

Here's a [companion video](#) that goes with this note.

Each output file is technically a valid CSV file, so you can also import the output in a spreadsheet app if you *really* want to crunch the numbers (or you really hate tables in terminal). If the app requires a `.csv` extension, you can use

```
$ cp tests/part-a/alu/student-output/alu-add-student.out student.csv
```

and import the resulting `.csv` file.

Debugging

Here's a [companion video](#) that goes with this section.

Similar to how you can step through your C code in GDB, you can also step through the test circuits in Logisim! For this example we'll be using the `alu-add` test.

Open `tests/part-a/alu/alu-add.circ` in Logisim. Among other things, one ROM (read-only memory) each feeds into the `Input_A`, `Input_B`, and `ALUSel` tunnels. These tunnels then feed into your ALU near the upper right. The ROMs contain corresponding values for the test being considered. Every clock cycle, the adder on the top left increments by one, which advances each ROM by one entry, thus feeding the next set of inputs to your ALU. If you tick the circuit a couple times (`File -> Tick Full Cycle` or the corresponding keyboard shortcut), you can see the test circuit advance through each set of inputs and your ALU's corresponding outputs. If you want to start over, use `Simulate -> Reset Simulation` (`Command/Control + R`).

Now, let's see what your ALU is actually doing with the inputs. Right-click your ALU, and select `View alu`. Your ALU circuit will appear, with the input values for the current test cycle already on the ALU input pins. With this, you can see exactly what your ALU is doing in every line from the output files! The `Poke Tool` will be very useful here.

Note: Edits to the test circuit, including the ALU we just inspected, **will not be saved**. Avoid making edits in the test circuit, as they may be lost!

Task 2: Register File (RegFile)

Fill in `regfile.circ` so that it contains 32 registers that can be written to and read from.

Input Name	Bit Width	Description
<code>rs1</code>	5	Determines which register's value is sent to the <code>Read_Data_1</code> output
<code>rs2</code>	5	Determines which register's value is sent to the <code>Read_Data_2</code> output
<code>rd</code>	5	The register to write to on the next rising edge of the clock (if <code>RegWEn</code> is 1)
<code>Write_Data</code>	32	The data to write into <code>rd</code> on the next rising edge of the clock (if <code>RegWEn</code> is

Input Name	Bit Width	Description
		1)
RegWEn	1	Determines whether data is written to the register file on the next rising edge of the clock
clk	1	Clock input

Output Name	Bit Width	Description
Read_Data_1	32	The value of the register identified by rs1
Read_Data_2	32	The value of the register identified by rs2
ra	32	The value of ra (x1)
sp	32	The value of sp (x2)
t0	32	The value of t0 (x5)
t1	32	The value of t1 (x6)
t2	32	The value of t2 (x7)
s0	32	The value of s0 (x8)
s1	32	The value of s1 (x9)
a0	32	The value of a0 (x10)

- The 8 constant output registers are included in the output of the **regfile** circuit for testing and debugging purposes. Make sure to connect these 8 output pins to their corresponding registers.
- The **x0** register should always contain the 0 value, even if an instruction tries writing to it.

Some additional tips:

- Take advantage of copy-paste! It might be a good idea to make one register completely and use it as a template for the others to avoid repetitive work. You can duplicate a selected component or group of components in Logisim using **Ctrl/Cmd + D**.
- The **Enable** pin on the built-in register may come in handy.

Testing

We've provided the autograder RegFile tests in the `tests/part-a/regfile/` directory. You can run these with:

```
$ python3 test.py tests/part-a/regfile/
```

Refer to the testing section of Task 1 for more information on test outputs.

Task 3: Immediate Generator

For the rest of Part A, we will be creating just enough of the CPU to execute the `addi` instruction. In Part B, you will revisit these circuits and expand them to support more instructions.

Fill in the immediate generator in `imm-gen.circ` (not the `imm_gen` subcircuit in `cpu.circ`) so that it can generate immediates for the `addi` instruction. You can ignore other immediate types for now.

Input Name	Bit Width	Description
<code>inst</code>	32	The instruction being executed
<code>ImmSel</code>	3	Value determining how to reconstruct the immediate (you can ignore this for now)
<code>imm</code>	32	Value of the immediate in the instruction (assume the instruction is <code>addi</code> for now)

Output Name	Bit Width	Description
<code>imm</code>	32	Value of the immediate in the instruction (assume the instruction is <code>addi</code> for now)

Task 4: Datapath

Fill in `cpu.circ` so that it contains a datapath for a single-cycle (not pipelined) processor that can execute the `addi` instruction.

Here are the inputs and outputs to the processor. You can leave most of them unchanged in this task, since they are not needed for the `addi` instruction.

Input Name	Bit Width	Description
READ_DATA	32	Data at WRITE_ADDRESS from memory
INSTRUCTION	32	The instruction at memory address PROGRAM_COUNTER
CLOCK	1	Clock input

Output Name	Bit Width	Description
ra	32	The value of ra (x1)
sp	32	The value of sp (x2)
t0	32	The value of t0 (x5)
t1	32	The value of t1 (x6)
t2	32	The value of t2 (x7)
s0	32	The value of s0 (x8)
s1	32	The value of s1 (x9)
a0	32	The value of a0 (x10)
tohost	32	The contents of CSR
WRITE_ADDRESS	32	The address in memory to read from or write to
WRITE_DATA	32	Data to write to memory
WRITE_ENABLE	4	The write enable mask for writing data to memory
PROGRAM_COUNTER	32	Address of the INSTRUCTION input

We know that trying to build a datapath from scratch might be intimidating, so the rest of this section offers more detailed guidance for creating your processor.

Recall the five stages for executing an instruction:

1. Instruction Fetch (IF)
2. Instruction Decode (ID)
3. Execute (EX)

4. Memory (MEM)
5. Write Back (WB)

Task 4.1: Instruction Fetch

We have already provided a simple implementation of the program counter. It is a 32-bit register that increments by 4 on each clock cycle. The `PROGRAM_COUNTER` is connected to IMEM (instruction memory), and the `INSTRUCTION` is returned from IMEM.

Nothing for you to implement in this sub-task!

Task 4.2: Instruction Decode

In this step, we need to break down the `INSTRUCTION` input and send the bits to the right subcircuits.

- ▶ What type of instruction is `addi`? What are the different fields in the instruction, and which bits correspond to each field?
- ▶ In Logisim, what tool would you use to split out different groups of bits?
- ▶ Which fields should connect to the register file? Which inputs of the register file should they connect to?
- ▶ What needs to be connected to the immediate generator?

Task 4.3: Execute

In this step, we will use the decoded instruction fields to compute the actual instruction.

- ▶ What two data values (`A` and `B`) should the `addi` instruction input to the ALU?
- ▶ What `ALUSel` value should the instruction input to the ALU?

Task 4.4: Memory

The `addi` instruction doesn't use memory, so there's nothing for you to implement in this sub-task!

The memory stage is where the memory can be written to using store instructions and read from using load instructions. Because the `addi` instruction does not use memory, we do not have to worry about it for Part A. Please ignore the DMEM and leave its I/O pins undriven.

Task 4.5: Write Back

In this step, we will write the result of our `addi` instruction back into a register.

► What data is the `addi` instruction writing, and where is the instruction writing this data to?

Testing

Here's a [companion video](#) that goes with this section.

Each CPU test is a copy of the `run.circ` file included with the starter code that has instructions loaded into its IMEM. When you run the `test.py` script, it runs Logisim in the background. The clock ticks, the program counter is incremented, and the values in each of the outputs is printed to a `.out` file in the `student-outputs` directory.

Let's take the single-cycle `cpu-addi-basic` sanity test as an example. It has 4 `addi` instructions (see `tests/part-a/addi/inputs/cpu-addi-basic.s`). Open `tests/part-a/addi/cpu-addi-basic.circ` in Logisim, and take a closer look at the various parts of the test file. At the top, you'll see the place where your CPU is connected to the test outputs. With the starter code, you'll see lots of `UUUU`s or `XXXX`s; when your CPU is working, this should not be the case. Your CPU takes in one input (`INSTRUCTION`), and along with the values in each of the registers, it has an additional output: `PROGRAM_COUNTER`, or the address of the instruction to be fetched from IMEM to be executed the next clock cycle.

As you can see, there are many specifically-positioned wires connected to specific input/output pins on your CPU. Make sure that you **do not** edit the provided input/output pins or add new ones, as this will change the shape of the CPU circuit, and as a result the connections in the test files may no longer work properly.

Below the CPU, you'll see instruction memory. The hex for the `addi` instructions has been loaded into instruction memory. Instruction memory takes in one input (called `PROGRAM_COUNTER`) and outputs the instruction at that address. `PROGRAM_COUNTER` is a 32-bit value, but because Logisim caps the size of ROM units at 2^{16} bytes, we have to use a splitter to get only 14 bits from `PROGRAM_COUNTER` (ignoring the bottommost two bits). Notice that `PROGRAM_COUNTER` is a **byte address**, not a word address.

So what happens when the clock ticks? Each tick of the clock increments an input in the test file called `Time_Step`. The clock will continue to tick until `Time_Step` is equal to the halting constant for that test file (for this particular test file, the halting constant is 6). At that point, the `test.py` script will print the values in each of the outputs to the respective `.out` file. Our tests will compare this output to the expected; if your output is different, you will fail the test.

We've provided the autograder tests for `addi` (Task 3) in the `tests/part-a/addi/` directory. You can run these with:

```
$ python3 test.py tests/part-a/addi/
```

Task 5: Part A README Update

Your last task for Part A is to fill in the `README.md`. Write down how you implemented your circuits and components for this part (including ALU and RegFile, since you used them for `addi`!), and explain the reasoning behind your design choices. There's no specific format or length requirement here, so feel free to get creative!

Submission and Grading

Submit your repository to the Project 3A assignment on Gradescope. The autograder tests for Part A are the same as the tests you are running locally. Part A is worth 20% of your overall Project 3 grade.

- ALU (7)
- RegFile (8)
- `addi` (5)

Total: 20 points

☐ Dark Mode



Part B: More Instructions

In this part, you will expand your CPU to support more instructions and pipelining.

You can implement the instructions in any order you want (see the appendix for a full list of instructions and subcircuit definitions). If you'd like a bit more guidance, this spec helps you build up your CPU with small groups of instructions at a time.

Task 6: I-type Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
addi rd, rs1, imm	I	0x13	0x0		$R[rd] \leftarrow R[rs1] + \text{imm}$
slli rd, rs1, imm			0x1	0x00	$R[rd] \leftarrow R[rs1] \ll \text{imm}$
slti rd, rs1, imm			0x2		$R[rd] \leftarrow (R[rs1] < \text{imm}) ? 1 : 0$
xori rd, rs1, imm			0x4		$R[rd] \leftarrow R[rs1] \wedge \text{imm}$
srli rd, rs1, imm			0x5	0x00	$R[rd] \leftarrow R[rs1] \gg \text{imm}$
srai rd, rs1, imm			0x5	0x20	$R[rd] \leftarrow R[rs1] \gg \text{imm}$
ori rd, rs1, imm			0x6		$R[rd] \leftarrow R[rs1] \vee \text{imm}$
andi rd, rs1, imm			0x7		$R[rd] \leftarrow R[rs1] \& \text{imm}$

Recall that you already implemented **addi** in Part A. These I-type instructions use the same datapath as **addi**, so your datapath already supports them!

Control Logic

As you add logic to support more instructions in the next few tasks, you will need to add control logic to enable the relevant datapath components depending on the instruction being executed.

The control logic should take the instruction bits and output all the control signals needed to execute that instruction. There are two main approaches to implementing this logic:

1. **Suggested: Read-only memory (ROM):** Hard-code a table of control signals. The relevant bits of the instruction are used to index into the table and find the control bits. See "Control Logic: Getting Started with ROM" for more information.

Pros: Fewer wires, more hard-coding (makes the control logic easier to debug).

Cons: ROM is common for CISC ISAs (like x86) but is less common for RISC ISAs (like RISC-V).

Newer to this project (and may have some bugs we haven't discovered yet).

2. **Alternate: Hard-wired control:** Use logic gates (e.g. AND/OR/NOT gates, MUXes) to calculate the control bits from the instruction.

Pros: Used in real-world RISC-V systems, used in all past semesters.

Cons: More wires (harder to debug). Creating boolean equations with logic gates to calculate control bits generally takes longer than hardcoding bits in the ROM.

Regardless of which approach you choose, you should modify `control_logic.circ` in each task to implement your control logic.

Here is a summary of the control signals you should implement:

Signal	Bit Width	Purpose
PCSel	1	Selects the ALU input for all SB-type instructions where the branch is taken (according to the branch comparator output). Selects the PC+4 input for all other instructions.
ImmSel	1	Selects the instruction format so the immediate generator can extract the immediate correctly.
RegWEn	1	1 if the instruction writes to a register, and 0 otherwise.
BrUn	1	1 if the branch instruction is unsigned, and 0 if the branch instruction is signed. Don't care for all other instructions.
ASel	1	Selects whether to send the data in <code>rs1</code> or the PC to the ALU.
BSel	1	Selects whether to send the data in <code>rs2</code> or the immediate to the ALU.
ALUSel	4	Selects the correct operation for the ALU.
MemRW	1	1 if the instruction writes to memory, and 0 otherwise.

WBSe1	2	Selects whether to write the memory read from DMEM, the ALU output, or PC+4 to rd.
-------	---	--

Control Logic: Getting Started with ROM

- `control-logic.circ` already has an empty ROM. To complete the control logic using ROM, you need to do the following:

► Details

► Details

► Details

► Details

- We do not suggest modifying the inputs or outputs to the existing starter `control_logic` circuit if you are using a ROM-based implementation, since the ROM and spreadsheet cannot be easily modified to output additional control signals. If you would like to add additional control signals, we suggest using hard-wired control logic to calculate them, like you did for PCSe1.

Task 6.1: Control Logic

Refer to the control logic section above for instructions.

Testing and Debugging

Before you begin, keep in mind that debugging and testing should be conducted as you are implementing your CPU. It is **A LOT** easier to debug one instruction type at a time versus doing them all together. If you prefer a video format, take a look at our videos on [creating tests](#) and [debugging](#). Please note that the debugging video uses an older version of the test generation script; refer to it only for debugging and use the "creating tests" videos as a guide on how to create tests.

Here are the same instructions in written form:

We've included a script (`tools/create-test.py`) that uses Venus to help you generate test circuits from RISC-V assembly! The process for writing custom tests is as follows:

1. Come up with a test, and write the RISC-V assembly instructions for that test, saving them in a file ending in `.s` in the `tests/part-b/custom/inputs/` folder. The name of this file will be the name of your test. Repeat if you have more tests.

- e.g. `tests/part-b/custom/inputs/sll-slli.s` and `tests/part-b/custom/inputs/beq.s`

Note: IMEM and DMEM are separate in Logisim, but combined in Venus. This means that if you write assembly code that tries to access memory overlapping with instructions, Venus will throw an error. Since counting exactly how many instructions your assembly code requires, and multiplying that by 4 can be annoying, we suggest you load/store using addresses greater than 0x3E8 (leaving space for 1000 bytes/250 instructions), and increase this offset if you have more instructions.

2. Generate the test circuits for your tests using `create-test.py`:

```
$ python3 tools/create-test.py tests/part-b/custom/inputs/sll-slli.s tests/part-b/custom/
```

Reminder: if you want to regenerate *everything*, you can take advantage of Bash globs:

```
$ python3 tools/create-test.py tests/part-b/custom/inputs/*.s
```

This should generate a couple new files to go with your assembly file:

```
tests/part-b/custom/:
- cpu-<TEST_NAME>.circ          # The new circuit for your test
- inputs/<TEST_NAME>.s          # The test file you wrote (uncha
- reference-outputs/cpu-<TEST_NAME>-pipelined-ref.out # The pipelined reference output
- reference-outputs/cpu-<TEST_NAME>-ref.out          # The single-cycle reference out
```

3. Now you can run the tests you just wrote!

```
$ python3 test.py tests/part-b/custom/sll-slli.circ tests/part-b/custom/beq.circ
```

Reminder: you can run all tests in a directory:

```
$ python3 test.py tests/part-b/custom/
```

By default, the number of cycles for a test will be just enough for all instructions in the test, as well as extra cycles for the register writeback and pipelining. If you wish to override this and simulate your code for a certain number of cycles, you can use the `--cycles` flag:

```
$ python3 tools/create-test.py --cycles 10 tests/part-b/custom/inputs/sll-slli.s
```

4. Each test will print out either:

```
PASSED test: <TEST_NAME> # Your student-output/<TEST_NAME>-student.out == reference-outp
FAILED test: <TEST_NAME> Your student-output/<TEST_NAME>-student.out != reference-output
```

5. If any of your tests fail, you will need to debug it more closely in Logisim

To debug any tests or see how the circuit is working for each instruction:

1. First find out where the output of your CPU implementation diverges from the staff solution. You can compare them using the `diff` command against the reference output and your output.

```
$ python3 tools/format-output.py tests/part-b/custom/reference-output/<TEST_NAME>-ref.out
$ python3 tools/format-output.py tests/part-b/custom/student-output/<TEST_NAME>-student.out
$ diff reference.out student.out
```

2. Open the generated .circ file using Logisim. You will now see that test harness and instruction ROM. You can now use that information from your diff to step through the circuit to where you know your circuit is malfunctioning.
3. Trace your wires and check outputs that are malfunctioning. Consider all possibilities:
 - If regA or regB is coming out wrong, maybe it was never set correctly
 - If DMEM read is coming out wrong, maybe it was never set correctly
 - If ALU output is wrong, maybe the control logic is malfunctioning
 - If the instruction binary is wrong, maybe branching was not implemented correctly
 - etc. etc., an important part of this project is to think about picking holes in your design and solving problems.

Task 6.2: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 7: R-type Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
-------------	------	--------	--------	------------------	-----------

add rd, rs1, rs2	R	0x33	0x0	0x00	$R[rd] \leftarrow R[rs1] + R[rs2]$
mul rd, rs1, rs2			0x0	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[31:0]$
sub rd, rs1, rs2			0x0	0x20	$R[rd] \leftarrow R[rs1] - R[rs2]$
sll rd, rs1, rs2			0x1	0x00	$R[rd] \leftarrow R[rs1] \ll R[rs2]$
mulh rd, rs1, rs2			0x1	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[63:32]$
mulhu rd, rs1, rs2			0x3	0x01	$(\text{unsigned}) R[rd] \leftarrow (R[rs1] * R[rs2])[63:32]$
slt rd, rs1, rs2			0x2	0x00	$R[rd] \leftarrow (R[rs1] < R[rs2]) ? 1 : 0$ (signed)
xor rd, rs1, rs2			0x4	0x00	$R[rd] \leftarrow R[rs1] \wedge R[rs2]$
srl rd, rs1, rs2			0x5	0x00	$(\text{unsigned}) R[rd] \leftarrow R[rs1] \gg R[rs2]$
sra rd, rs1, rs2			0x5	0x20	$(\text{signed}) R[rd] \leftarrow R[rs1] \gg R[rs2]$
or rd, rs1, rs2			0x6	0x00	$R[rd] \leftarrow R[rs1] R[rs2]$
and rd, rs1, rs2			0x7	0x00	$R[rd] \leftarrow R[rs1] \& R[rs2]$

Task 7.1: Datapath

Modify your datapath in `cpu.circ` so that it can support R-type instructions.

If you're stuck, read further for some guiding questions. As with Task 4, it may help to think about each of the five stages for executing an instruction.

► Instruction Fetch: How do R-type instructions affect the program counter?

► Instruction Decode: What do we need to read from the register file?

► Execute: What two data values (A and B) should an R-type instruction input to the ALU?

► Memory: Do R-type instructions write to memory?

► Write back: What data is the R-type instruction writing, and where is the instruction writing this data to?

Task 7.2: Control Logic

Refer to the control logic section above for instructions.

Task 7.3: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 8: SB-type Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
beq rs1, rs2, offset	SB	0x63	0x0		if(R[rs1] == R[rs2]) $PC \leftarrow PC + \{\text{offset}, 1b0\}$
bne rs1, rs2, offset			0x1		if(R[rs1] != R[rs2]) $PC \leftarrow PC + \{\text{offset}, 1b0\}$
blt rs1, rs2, offset			0x4		if(R[rs1] < R[rs2] (signed)) $PC \leftarrow PC + \{\text{offset}, 1b0\}$
bge rs1, rs2, offset			0x5		if(R[rs1] >= R[rs2] (signed)) $PC \leftarrow PC + \{\text{offset}, 1b0\}$
bltu rs1, rs2, offset			0x6		if(R[rs1] < R[rs2] (unsigned)) $PC \leftarrow PC + \{\text{offset}, 1b0\}$
bgeu rs1, rs2, offset			0x7		if(R[rs1] >= R[rs2] (unsigned)) $PC \leftarrow PC + \{\text{offset}, 1b0\}$

Task 8.1: Branch Comparator

Fill in the branch comparator subcircuit in [branch-comp.circ](#). This subcircuit takes two inputs and outputs the result of comparing the two inputs. We will use the output later for implementing

branches.

Signal Name	Direction	Bit Width	Description
rs1	Input	32	First value to compare
rs2	Input	32	Second value to compare
BrUn	Input	1	1 when an unsigned comparison is wanted, and 0 when a signed comparison is wanted
BrEq	Output	1	Set to 1 if the two values are equal
BrLt	Output	1	Set to 1 if the value in rs1 is less than the value in rs2

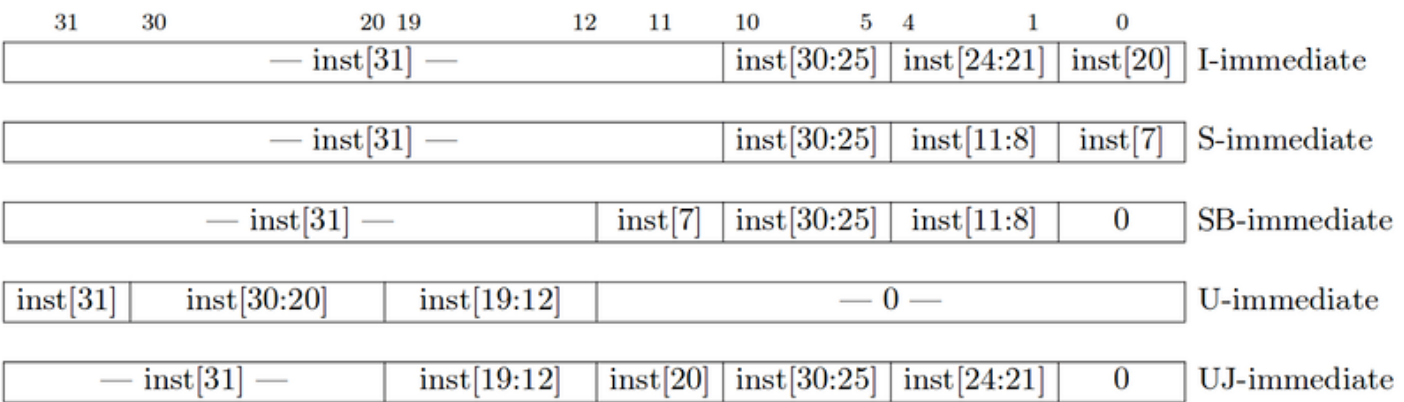
Task 8.2: Immediate Generator

Edit the immediate generator in `imm-gen.circ` so that it can generate immediates for SB-type instructions in addition to immediates for I-type instructions (which you implemented in Part A).

Recall that the bits of the immediate are stored in different bits of the instruction, depending on the type of the instruction. The `ImmSel` signal, which you will implement in the control logic, will determine which type of immediate this subcircuit should generate.

Also, recall that all immediates are 32 bits and sign-extended.

The immediate storage formats are listed below:



Task 8.3: Datapath

Modify your datapath in `cpu.circ` so that it can support SB-type instructions.

If you're stuck, read further for some guiding questions. As with Task 4, it may help to think about each of the five stages for executing an instruction.

- ▶ Instruction Fetch: How do SB-type instructions affect the program counter?
- ▶ Instruction Decode: What do we need to read from the register file?
- ▶ Execute: What two data values (A and B) should an SB-type instruction input to the ALU?
- ▶ Memory: Do SB-type instructions write to memory?
- ▶ Write back: What data is the SB-type instruction writing, and where is the instruction writing this data to?

Task 8.4: Control Logic

Refer to the control logic section above for instructions.

Task 8.5: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 9: Loading and Storing

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
lb rd, offset(rs1)	I	0x03	0x0		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset, byte}))$
lh rd, offset(rs1)			0x1		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset, half}))$
lw rd, offset(rs1)			0x2		$R[rd] \leftarrow \text{Mem}(R[rs1] + \text{offset, word})$

sb rs2, offset(rs1)	S	0x23	0x0		Mem(R[rs1] + offset) \leftarrow R[rs2] [7:0]
sh rs2, offset(rs1)			0x1		Mem(R[rs1] + offset) \leftarrow R[rs2] [15:0]
sw rs2, offset(rs1)			0x2		Mem(R[rs1] + offset) \leftarrow R[rs2]

Task 9.1: Immediate Generator

Edit the immediate generator in `imm-gen.circ` so that it can generate immediates for S-type instructions in addition to all the instruction types from previous tasks.

Recall that the bits of the immediate are stored in different bits of the instruction, depending on the type of the instruction. The `ImmSel` signal, which you will implement in the control logic, will determine which type of immediate this subcircuit should generate.

Also, recall that all immediates are 32 bits and sign-extended.

The immediate storage formats are listed below:

31	30	20	19	12	11	10	5	4	1	0	
— inst[31] —						inst[30:25]	inst[24:21]	inst[20]	I-immediate		
— inst[31] —						inst[30:25]	inst[11:8]	inst[7]	S-immediate		
— inst[31] —					inst[7]	inst[30:25]	inst[11:8]	0	SB-immediate		
inst[31]	inst[30:20]		inst[19:12]		— 0 —						U-immediate
— inst[31] —				inst[19:12]		inst[20]	inst[30:25]	inst[24:21]	0	UJ-immediate	

Conceptual Overview: Memory Unit

Since these instructions load from memory and store to memory, you will need to interact with the memory subcircuit (already implemented for you).

Signal Name	Direction	Bit Width	Description
WriteAddr	Input	32	Address to read/write to in Memory

Signal Name	Direction	Bit Width	Description
WriteData	Input	32	Value to be written to Memory
Write_En	Input	4	The write mask for instructions that write to Memory and zero otherwise
CLK	Input	1	Driven by the clock input to the CPU
ReadData	Output	32	Value of the data stored at the specified address

Due to Logisim size limitations, the memory unit only uses the lower 16 bits of the provided address, discarding the upper 16 bits. This means that the memory can only store 2^{16} bytes of data. The provided tests will always set the upper 16 bits of addresses to 0, and your tests should avoid using the upper 16 bits when interacting with memory.

The memory unit will always zero out the bottom two bits of the provided address and read 4 bytes starting at this modified address. In other words, the provided address will be rounded down to the nearest multiple of 4, and 4 bytes will be read in that range of 4 bytes.

For example, if the input address `WriteAddr` is `0x00001007`, the bottom two bits will be zeroed out to make `0x00001004`, and the output `ReadData` will be the 4 bytes at addresses `0x00001004`, `0x00001005`, `0x00001006`, and `0x00001007`.

The same process of zeroing out the bottom two bits applies for writing data as well, but the 4-bit `Write_En` input lets you select which of the 4 bytes in the `WriteData` input are written to the 4 bytes at the zeroed-out `WriteAddr` address. Each bit of this write mask enables writing to the corresponding byte of the word.

For example, if the input address `WriteAddr` is `0x00001007`, the write mask `Write_En` is `0b0010`, and the `WriteData` input is `0x11223344`, then the byte `0x33` will be written to `0x00001006`. If `Write_En` was changed to `0b0011`, then the byte `0x44` would also be written to `0x00001007`.

Conceptual Overview: Alignment

In this project, all memory accesses will be *aligned*. This means that a single load or store instruction will never cross a word boundary in memory.

All `lw` and `sw` instructions will use memory addresses that end in `0b00` (accessing `0b00`, `0b01`, `0b02`, and `0b03` in memory).

All **lh** and **sh** instructions will use memory addresses that end in either **0b00** (accessing **0b00** and **0b01**) or **0b10** (accessing **0b10** and **0b11**).

You should not implement any unaligned memory accesses in this project.

Task 9.2: Datapath

Modify your datapath in **cpu.circ** so that it can support loads and stores.

You should provide an address input **WriteAddr** to DMEM. Remember that the ALU calculates this address by adding the address in **rs1** and the offset immediate.

For load instructions, you should also add functionality in the write-back stage so that the data outputted by DMEM (and processed by your logic in Task 9.3) is written back to the **rd** register.

Task 9.3: Load

Remember that the address input **WriteAddr** will have its bottom two bits zeroed out. The DMEM will then output **ReadData**, which contains 4 bytes starting at this modified address. Implement logic in **cpu.circ** to extract and, if needed, sign-extend the relevant byte(s) from **ReadData**. Only the relevant byte(s) with appropriate sign-extending should be written back to the **rd** register.

For completeness, a table of scenarios you need to handle is provided below:

Instruction	Type	Opcode	Funct3	Bottom 2 bits of address	Value to put in rd
lb rd, offset(rs1)	I	0x03	0x0	0b00	SignExt(ReadData[7:0])
				0b01	SignExt(ReadData[15:8])
				0b10	SignExt(ReadData[23:16])
				0b11	SignExt(ReadData[31:24])
lh rd, offset(rs1)			0x1	0b00	SignExt(ReadData[15:0])
				0b10	SignExt(ReadData[31:16])
lw rd, offset(rs1)			0x2	0b00	ReadData

Task 9.3: Store

Implement logic in `cpu.circ` to set up the `WriteData` input and `Write_En` write mask to correctly store the data in `rs2` to memory.

For completeness, a table of scenarios you need to handle is provided below:

Instruction	Type	Opcode	Funct3	Bottom 2 bits of address	WriteData	Write_En
sb rs2, offset(rs1)	S	0x23	0x0	0b00	{24'b0, R[rs2] [7:0]}	0001
				0b01	{16'b0, R[rs2] [7:0], 8'b0}	0010
				0b10	{8'b0, R[rs2][7:0], 16'b0}	0100
				0b11	{R[rs2][7:0], 24'b0}	1000
sh rs2, offset(rs1)			0x1	0b00	{16'b0, R[rs2] [15:0]}	0011
				0b10	{R[rs2][15:0], 16'b0}	1100
sw rs2, offset(rs1)			0x2	0b00	R[rs2]	1111

Note that for any non-store instruction (i.e. when your `MemRW` control signal is 0), the `Write_En` write mask should be set to `0000`.

Task 9.4: Control Logic

Refer to the control logic section above for instructions.

Task 9.5: Testing and Debugging

Refer to the testing and debugging section above for instructions.

Task 10: All Other Instructions

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
auipc rd, offset	U	0x17			$R[rd] \leftarrow PC + \{\text{offset}, 12b0\}$
lui rd, offset		0x37			$R[rd] \leftarrow \{\text{offset}, 12b0\}$
jal rd, imm	UJ	0x6f			$R[rd] \leftarrow PC + 4$ $PC \leftarrow PC + \{\text{imm}, 1b0\}$
jalr rd, rs1, imm	I	0x67	0x0		$R[rd] \leftarrow PC + 4$ $PC \leftarrow R[rs1] + \{\text{imm}\}$

Task 10.1: Datapath

Modify your datapath in `cpu.circ` so that it can support these instructions. Most of these instructions are already supported by your datapath so far.

To support `jalr`, you should connect $PC+4$ to your multiplexer in the write-back stage so that $PC+4$ can be written back to `rd`.

Task 10.2: Control Logic

Refer to the control logic section above for instructions.

Task 10.3: Testing and Debugging


Refer to the testing and debugging section above for instructions.

Task 11: CSRW

The instructions you need to implement for this task are listed below:

Instruction	Type	Opcode	Funct3	Funct7/Immediate	Operation
csrw rd, csr, rs1	I	0x73	0x1		$CSR[csr] \leftarrow R[rs1]$
csrw rd, csr, uimm			0x5		$CSR[csr] \leftarrow \{\text{uimm}\}$

The instruction format for these instructions is shown below:



Conceptual Overview: Control Status Registers (CSRs)

Control Status Registers (CSRs) are used to hold additional information about the instructions being executed, such as debugging information or testing flags. These are unrelated to the registers and memory from the previous tasks.

To fully test your implementation, the project autograder requires one CSR (`tohost = 0x51E`). This means that for this project, the `csr` field will always be `0x51E` in this project.

Also, for this project, `rd` will always be 0, so you don't need to worry about CSR instructions writing to a register.

Task 11.1: Immediate Generator

Edit the immediate generator in `imm-gen.circ` so that it can generate immediates for the `csrwi` instruction (in addition to all the previous immediates you've implemented).

Note that the immediate is in bits 19-15 (labeled `uimm[4:0]`) and should be zero-extended, not sign-extended.

Task 11.2: Datapath

Modify your datapath in `cpu.circ` so that it can support these instructions.

You will need to interact with the CSR subcircuit (already implemented for you).

Signal Name	Direction	Bit Width	Description
CSR_address	Input	12	Input CSR register address
CSR_din	Input	32	Value to write into specified CSR register
CSR_WEn	Input	1	Write enable (from control logic)
clk	Input	1	Clock input
tohost	Output	32	Output of the <code>tohost</code> register

You should connect `CSR_address` to the 12-bit `csr` field in the instruction.

For `csrw` instructions, `CSR_din` should be the data in `rs1`. For `csrwi` instructions, `CSR_din` should be the immediate from the immediate generator.

Task 11.3: Control Logic

Refer to the control logic section above for instructions.

Note that there are two additional control signals for the CSR instructions:

Signal	Bit Width	Purpose
<code>CSRSe1</code>	1	Selects whether the instruction is <code>csrw</code> or <code>csrwi</code> . Don't care for all other instructions.
<code>CSRWen</code>	1	Selects whether the instruction is a CSR instruction. 1 for <code>csrw</code> and <code>csrwi</code> , 0 otherwise.

Task 12: Pipelining

In this task, you will implement a 2-stage pipeline in your CPU:

1. **Instruction Fetch:** An instruction is fetched from the instruction memory.
2. **Execute:** The instruction is decoded, executed, and committed (written back). This is a combination of the remaining four stages of a classic five-stage RISC-V pipeline (ID, EX, MEM and WB).

Some things to consider:

- Will the IF and EX stages have the same or different `PC` values?
- Do you need to store the `PC` between the pipelining stages?
- What hazards are present in this two-stage pipeline?

Note: During the first cycle, the instruction register sitting between the pipeline stages won't contain an instruction loaded from memory. What should the second stage do? Luckily, Logisim automatically sets registers to zero on reset, so the instruction pipeline register will automatically start with a `nop`! If you wish, you can depend on this behavior of Logisim.

Since your CPU will support branch and jump instructions, you'll need to handle control hazards that occur when branching.

- The instruction immediately after a branch or jump should not be executed if a branch is taken. This makes your task a bit more complex. By the time you have figured out that a branch or jump is in the execute stage, you have already accessed the instruction memory and pulled out (possibly) the wrong instruction. Therefore, you will need to flush the instruction that is being fetched (next instruction) if the instruction under execution (current instruction) is a taken branch or jump.
- Instruction flushing for this project must be accomplished by MUXing a `nop` into the instruction stream and sending the `nop` into the Execute stage instead of using the fetched instruction. You can use `addi x0, x0, 0 (0x00000013)`, for this purpose; other `nop` instructions will work too. You should flush only if a branch is taken (do not flush if it is not taken). You should always flush the next instruction when jumping.
- Note: you should not solve this issue by calculating branch offsets in the IF stage. We compare your output against a reference output that uses `nop`, so a solution that doesn't use `nop` may not match the reference output even if it is a conceptually correct solution.

Some more things to consider:

- To MUX a `nop` into the instruction stream, do you place it *before* or *after* the instruction register?
- What address should be requested next while the EX stage executes a `nop`? Is this different than normal?

Testing

We've provided some basic sanity tests for your pipelined CPU in the `tests/part-b/sanity/` directory (same tests as in Task 6). You can run these with:

```
$ python3 test.py --pipelined tests/part-b/sanity/
```

Note: since your CPU is pipelined at this point, you need to run the pipelined tests using the `--pipelined` (or `-p`) flag. If you run the single-cycle tests (i.e. omit the `--pipelined` flag) after pipelining your CPU, your CPU should now fail those tests! Think about why this happens...

Similarly, you can also run the pipelined version of your custom tests:

```
$ python3 test.py --pipelined tests/part-b/custom/
```

Note: because you're implementing a 2-stage pipelined processor and the first instruction writes on the rising edge of the second clock cycle, the effects of your instructions will have a 2 instruction delay. For example, let's look at the first instruction of `tests/part-b/sanity/inputs/addi.s`, `addi t0, x0, -1`. If you inspect the pipelined reference output (`tests/part-b/sanity/reference-output/cpu-addi-pipelined-ref.out`), you'll see that `t0` doesn't show changes until the third cycle.

Refer to the section from Project 3A for more info on using these tests. Keep in mind that you're working with a pipelined circuit from this task onward.

Task 13: Part B README Update

Time to update your `README.md`! Once again, write down how you implemented your circuits and components for this part (including the various subcircuits you used), and explain the reasoning behind your design choices. In particular, we want to see:

- How you designed your control logic
- Advantages/Disadvantages of your design
- Best/Worst bug or design challenge you encountered, and your solution to it

Your additions to the README should be at least 512 characters (although something more than the bare minimum would be nice), but other than that feel free to get creative!

Submission and Grading

Submit your assignment to the Project 3B submission on Gradescope. Part B is worth 80% of your overall Project 3 grade.

- Unit Tests (35%)
- Test Coverage (10%)
- Integration (5%)
- Edge (29%)
- README (1%)



Appendix: Logisim Tips

This appendix contains some helpful Logisim tips and pitfalls to avoid.

Wiring

- If you want to know more details about each component, go to [Help -> Library Reference](#) for more information on the component and its inputs and outputs.
- Use tunnels! They will make your wiring cleaner and easier to follow, and will reduce your chances of encountering crossed wires or unexpected errors.
- Ensure you name your tunnels correctly. The labels are case sensitive!
- You can hover your cursor over an input/output on a component to get slightly more information about that input/output.

Wiring Pitfalls

- Your circuits should always fit in the provided harnesses. This means that you should not edit the provided input/output pins or add new ones. To ensure your circuit fits into the harness, you can open the harnesses in the [harnesses](#) folder and check that there are no errors.
- Don't create new [.circ](#) files. You can make additional subcircuits if you want, but they must be in existing files.

Subcircuits

- Note that if you modify a subcircuit, and another circuit file uses that subcircuit, you will need to close and re-open the outer circuit to load the changes from the subcircuit. For example, if you modify [imm-gen.circ](#), you should close and re-open [cpu.circ](#) to load your changes.
- When modifying a subcircuit, you should always open up the subcircuit file. For example, you should modify [imm-gen.circ](#), not the [imm-gen](#) subcircuit in [cpu.circ](#).

Signal Tips

- The clock input signal ([clk](#)) can be sent into subcircuits or attached directly to the clock inputs of memory units in Logisim, but should not otherwise be gated (i.e., do not invert it, do not [AND](#) it with anything, etc.).

- We recommend not using the **Enable** input on your MUXes. In fact, you can turn that attribute off (**Include Enable?**). We also recommend that you disable the **Three-state?** attribute (if the plexer has it).

☐ Dark Mode